


AI as the Design of Agents

**= a unifying view for the bag of
techniques that AI encompasses**

Tuomas Sandholm
Carnegie Mellon University
Computer Science Department

An agent and its environment



How to design an intelligent agent?

- **Definition:** An *agent* perceives its environment via sensors and acts in that environment with its effectors.
Hence, an agent gets percepts one at a time, and maps this percept sequence to actions (one action at a time)
- **Properties:**
 - Autonomous
 - Interacts with other agents plus the environment
 - Reactive to the environment
 - Pro-active (goal-directed)

Examples of agents in different types of applications

Agent type	Percepts	Actions	Goals	Environment
Medical diagnosis system	Symptoms, findings, patient's answers	Questions, tests, treatments	Healthy patients, minimize costs	Patient, hospital
Satellite image analysis system	Pixels of varying intensity, color	Print a categorization of scene	Correct categorization	Images from orbiting satellite
Part-picking robot	Pixels of varying intensity	Pick up parts and sort into bins	Place parts in correct bins	Conveyor belts with parts
Refinery controller	Temperature, pressure readings	Open, close valves; adjust temperature	Maximize purity, yield, safety	Refinery
Interactive English tutor	Typed words	Print exercises, suggestions, corrections	Maximize student's score on test	Set of students


Definition of ideal rational agent

Ideal Rational Agent: For each possible percept sequence, such an agent does whatever action is expected to maximize its performance measure, on the basis of the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

What do you think? Is this an acceptable definition?

Not looking left when crossing the street: If I don't see a car coming from the left, it is rational to cross the street?

No. Should also consider taking *information gathering actions*.



Agent's strategy

- Agent's *strategy* is a mapping from percept sequence to action
- How to encode an agent's strategy?
 - Long list of what should be done for each possible percept sequence
 - vs. shorter specification (e.g. algorithm)

WARNING:

Might not get what you ask for in the performance measure

- Cleaning robot
 - Pick up as much trash as possible
- Vehicle route optimization
 - Maximize utilizations => Driving fully loaded
 - Capitalizing on oddities in tariff list => Renegotiation
 - Don't include solution method in the criterion

agent = architecture + **program**

This course concentrates on the program

Physical agents vs. software agents
(software agents = softbots)

Skeleton agent

```
function SKELETON-AGENT (percept) returns action
  static: memory, the agent's memory of the world

  memory ← UPDATE-MEMORY(memory,percept)
  action ← CHOOSE-BEST-ACTION(memory)
  memory ← UPDATE-MEMORY(memory, action)
  return action
```

On each invocation, the agent's memory is updated to reflect the new percept, the best action is chosen, and the fact that the action was taken is also stored in the memory. The memory persists from one invocation to the next.

Input = Percept, not history

NOTE: Performance measure is not part of the agent

Examples of how the agent function can be implemented

More
sophisticated



1. Table-driven agent
2. Simple reflex agent
3. Reflex agent with internal state
4. Agent with explicit goals
5. Utility-based agent

1. Table-driven agent

```
function TABLE-DRIVEN-AGENT (percept) returns action
  static: percepts, a sequence, initially empty
         table, a table, indexed by percept sequences, initially fully specified

  append percept to the end of percepts
  action  $\leftarrow$  LOOKUP(percepts, table)
  return action
```

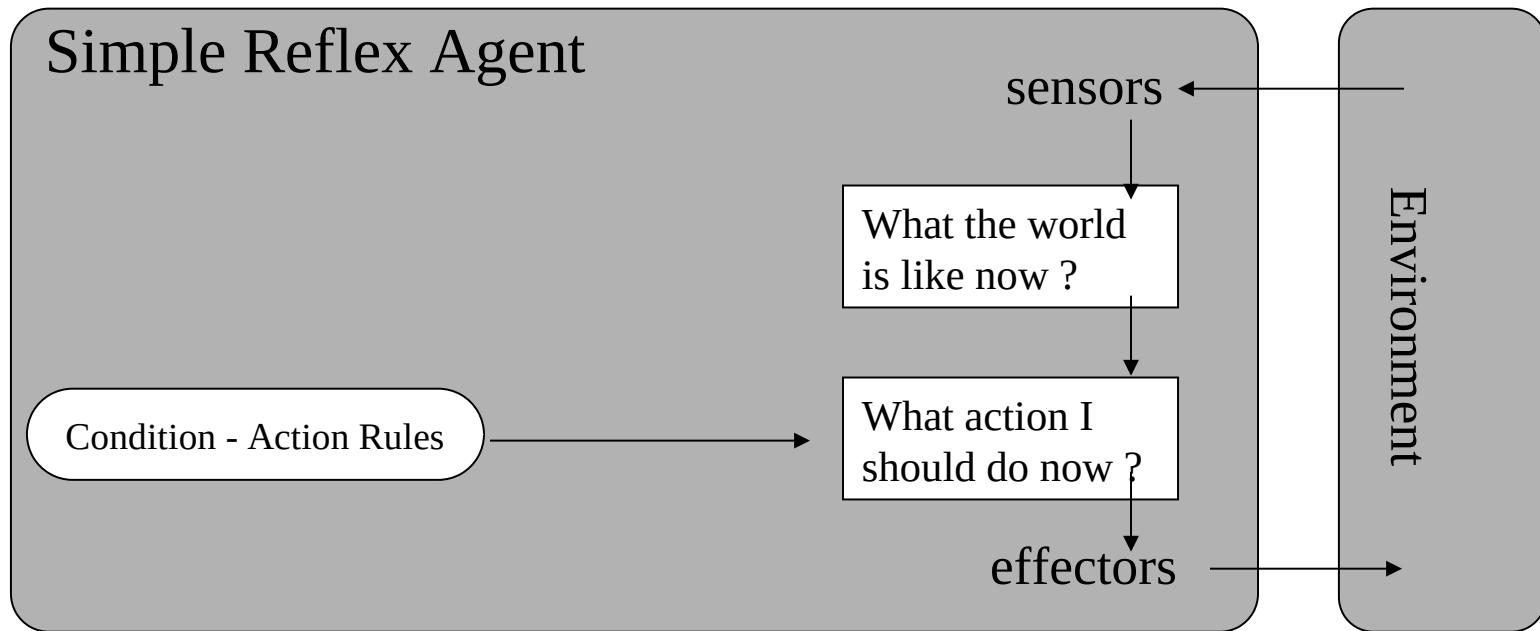
An agent based on a prespecified lookup table. It keeps track of percept sequence and just looks up the best action

- **Problems**

- Huge number of possible percepts (consider an automated taxi with a camera as the sensor) => lookup table would be huge
- Takes long time to build the table
- Not adaptive to changes in the environment; requires entire table to be updated if changes occur

2. Simple reflex agent

- Differs from the lookup table based agent in that the condition action rule (that determines the action) is already higher-level interpretation of the percepts
 - Percepts could be e.g. the pixels on the camera of the automated taxi



```
function SIMPLE-REFLEX-AGENT(percept) returns action
  static: rules, a set of condition-action rules

  state ← INTERPRET-INPUT (percept)
  rule ← RULE-MATCH (state,rules)
  action ← RULE-ACTION [rule]
  return action
```

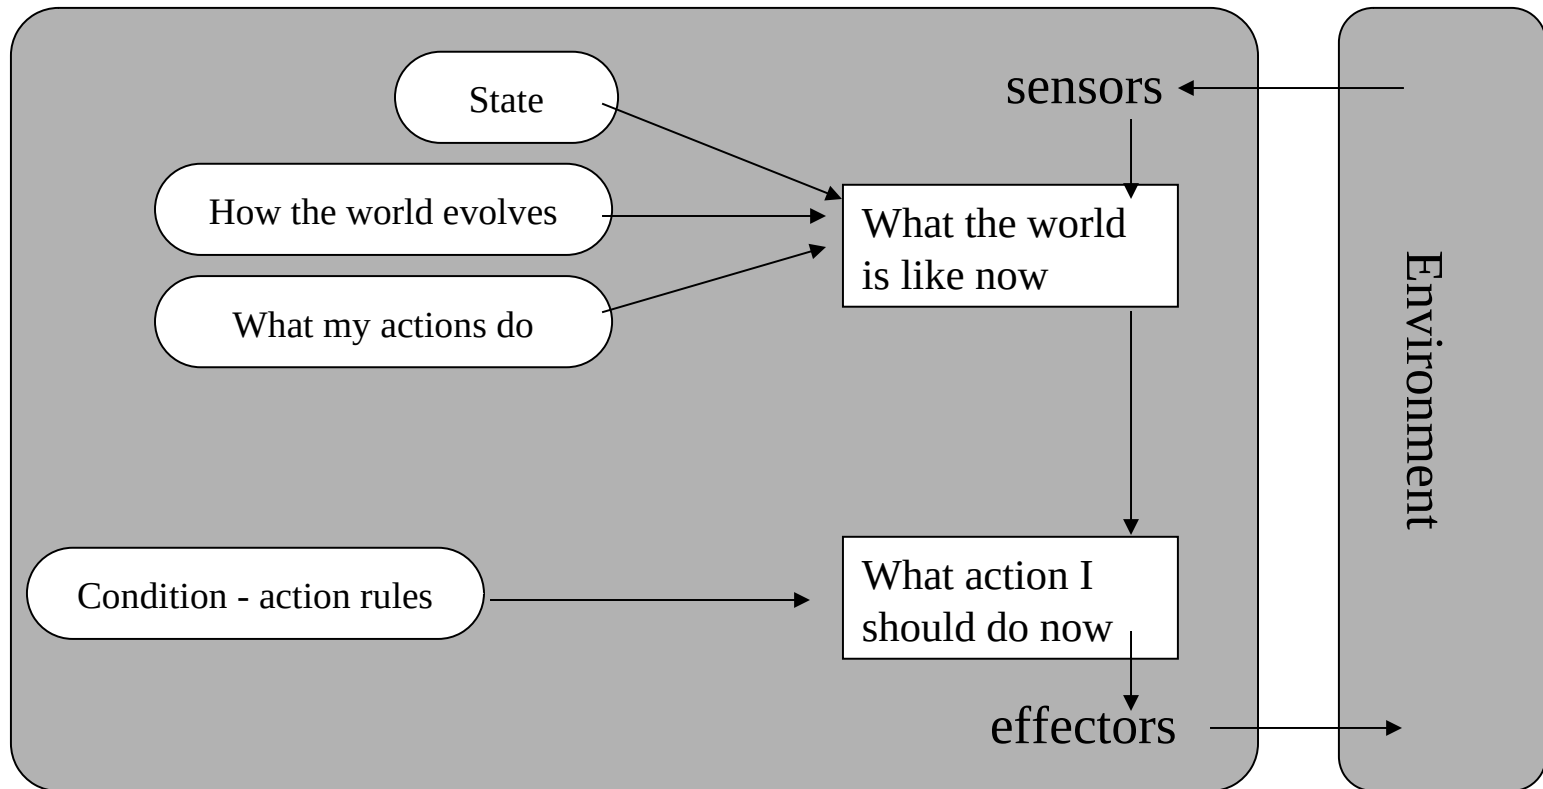
First match.
No further matches sought.
Only one level of deduction.

A simple reflex agent works by finding a rule whose condition matches the current situation (as defined by the percept) and then doing the action associated with that rule.

Simple reflex agent...

- **Table lookup of condition-action pairs defining all possible condition-action rules necessary to interact in an environment**
 - e.g. if car-in-front-is-breaking then initiate breaking
- **Problems**
 - Table is still too big to generate and to store (e.g. taxi)
 - Takes long time to build the table
 - No knowledge of non-perceptual parts of the current state
 - Not adaptive to changes in the environment; requires entire table to be updated if changes occur
 - Looping: Can't make actions conditional

3. Reflex agent *with internal state*



Reflex agent with internal state ...

```
function REFLEX-AGENT-WITH-STATE (percept) returns action
  static: state, a description of the current world state
         rules, a set of condition-action rules

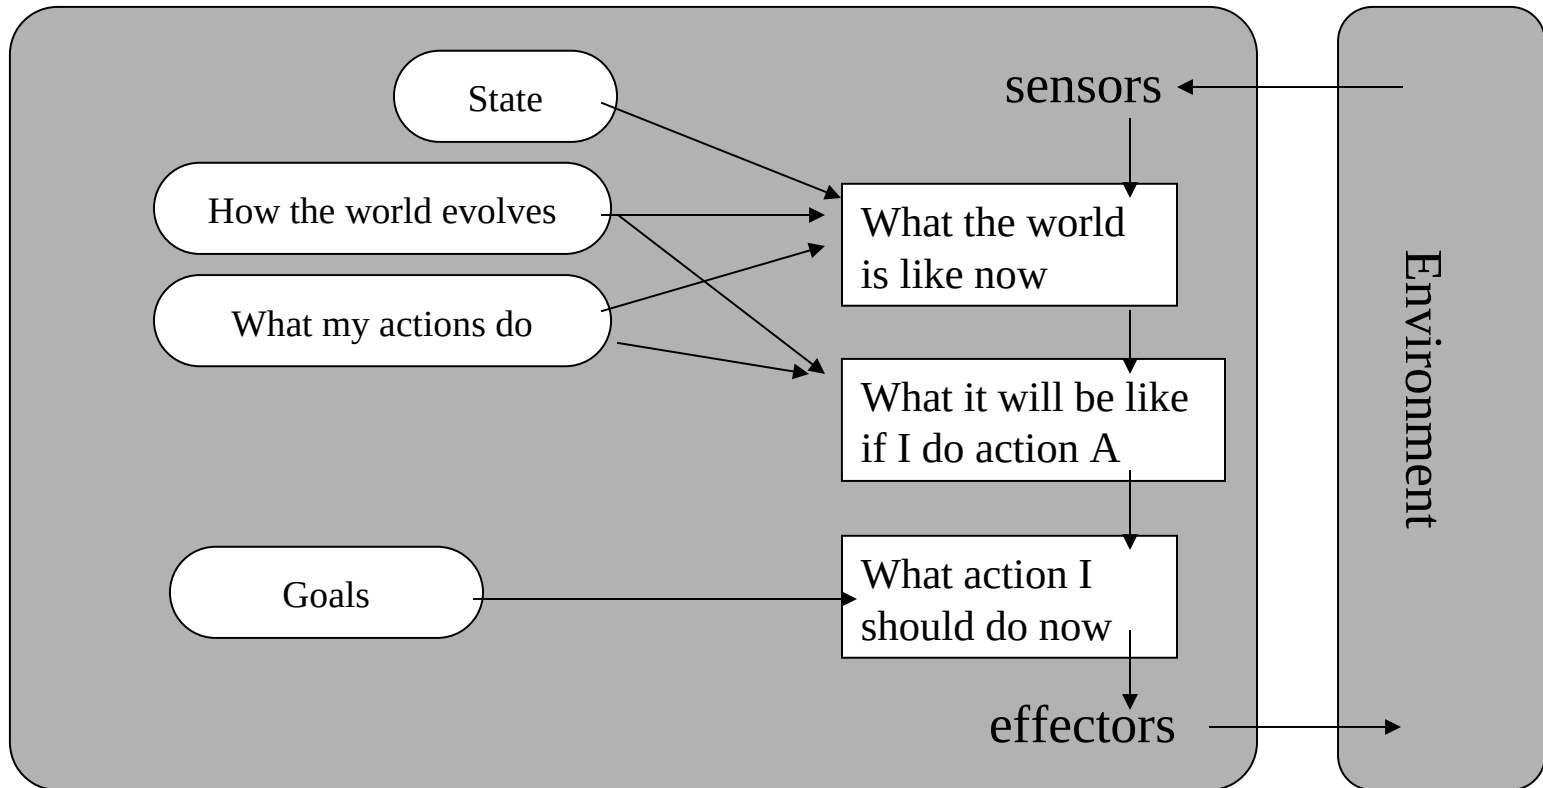
  state  $\leftarrow$  UPDATE-STATE (state, percept)
  rule  $\leftarrow$  RULE-MATCH (state, rules)
  action  $\leftarrow$  RULE-ACTION [rule]
  state  $\leftarrow$  UPDATE-STATE (state, action)
  return action
```

A reflex agent with internal state works by finding a rule whose condition matches the current situation (as defined by the percept and the stored internal state) and then doing the action associated with that rule.

Reflex agent with internal state ...

- Encode “internal state of the world to remember the past as contained in earlier percepts
- Needed because sensors do not usually give the entire state of the world at each input, so perception of the environment is captured over time. “State” used to encode different “world states” that generate the same immediate percept
- Requires ability to represent change in the world with/without the agent
 - one possibility is to represent just the latest state, but then cannot reason about hypothetical courses of action

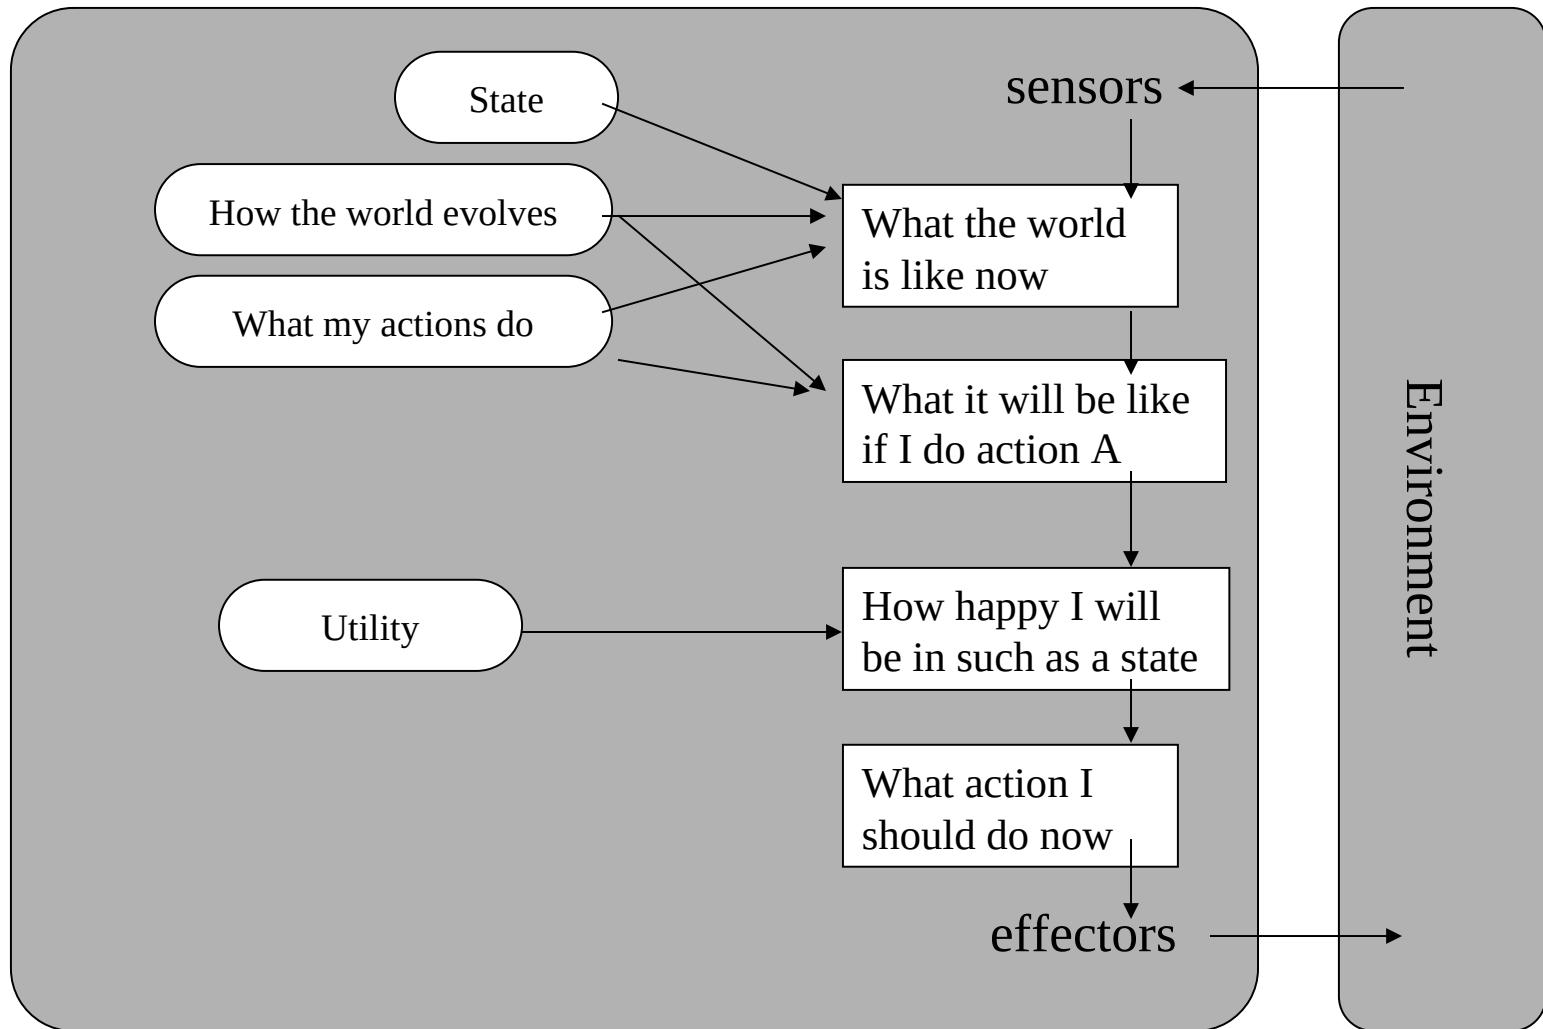
4. Agent with explicit goals



Agent with explicit goals ...

- Choose actions so as to achieve a (given or computed) goal = a description of desirable situations. e.g. where the taxi wants to go
- Keeping track of the current state is often not enough – need to add goals to decide which situations are good
- Deliberative instead of reactive
- May have to consider long sequences of possible actions before deciding if goal is achieved – involves considerations of the future, “what will happen if I do...?” (search and planning)
- More flexible than reflex agent. (e.g. rain / new destination) In the reflex agent, the entire database of rules would have to be rewritten

5. Utility-based agent



Utility-based agent ...

- When there are multiple possible alternatives, how to decide which one is best?
- A goal specifies a crude destination between a happy and unhappy state, but often need a more general performance measure that describes “degree of happiness”
- Utility function U : State \rightarrow Reals indicating a measure of success or happiness when at a given state
- Allows decisions comparing choice between conflicting goals, and choice between likelihood of success and importance of goal (if achievement is uncertain)

Properties of environments

Properties of environments

Accessible (observable) : The agent's sensory apparatus gives it access to the complete state of the environment

Deterministic: The next state of the environment is completely determined by the current state and the actions selected by the agent

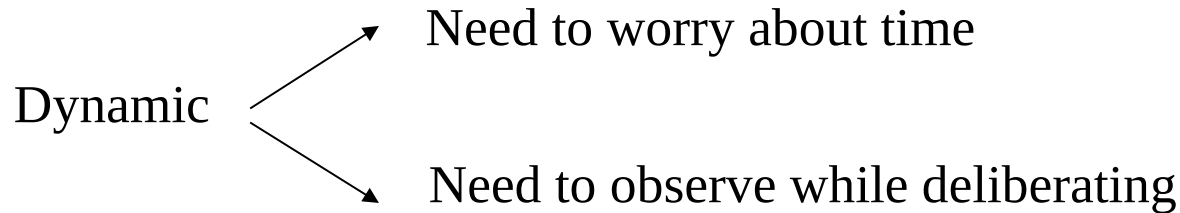
Subjective non-determinism

- Limited memory (poker)
- Too complex environment to model directly (weather, dice)
- Inaccessibility

Episodic: The agent's experience is divided into independent “episodes,” each episode consisting of agent perceiving and then acting. Quality of action depends just on the episode itself, because subsequent episodes do not depend on what actions occur in previous episodes. → Do not need to think ahead

Properties of environments ...

Static: If the environment can change while the agent is deliberating, then the environment is dynamic; otherwise it's static.



Discrete: There are a limited number of distinct, clearly defined percepts and actions we say that environment is discrete.

Environment	Accessible	Deterministic	Episodic	Static	Discrete
Chess with a clock	Yes	Yes	No	Semi	Yes
Chess without a clock	Yes	Yes	No	Yes	Yes
Poker	No	No	No	Yes	Yes
Backgammon	Yes	No	No	Yes	Yes
Taxi driving	No	No	No	No	No
Medical diagnosis system	No	No	No	No	No
Image-analysis system	Yes	Yes	Yes	Semi	No
Part-picking robot	No	No	Yes	No	No
Refinery controller	No	No	No	No	No
Interactive English tutor	No	No	No	No	Yes

Running the agents and the environment

procedure RUN-ENVIRONMENT (*state*, UPDATE-FN, *agents*, *termination*)

inputs: *state*, the initial state of the environment

UPDATE-FN, function to modify the environment

agents, a set of agents

termination, a predicate to test when we are done

repeat

for each *agent* **in** *agents* **do**

PERCEPT[*agent*] \leftarrow GET-PERCEPT(*agent*,*state*)

end

for each *agent* **in** *agents* **do**

ACTION[*agent*] \leftarrow PROGRAM[*agent*] (PERCEPT[*agent*])

end

state \leftarrow UPDATE-FN(*actions*, *agents*, *state*)

until *termination* (*state*)